

PeerSim HOWTO: build a new protocol for the peersim simulation framework

Gian Paolo Jesi (jesi@cs.unibo.it)

September 26, 2005

1 Introduction

NOTE: This tutorial revision covers peersim release 0.5 topics.

NOTE: if the reader is not a peersim newbie, she may be interested in which major changes have been introduced; a small, but fast reference is provided in Appendix C.

This tutorial is aimed to give you a step by step guide to build from scratch a new peersim application (<http://sourceforge.net/projects/peersim>): a framework to experiment with large scale P2P overlay networks. In this tutorial it supposed that you and/or your workstation have:

- knowledge of O.O. programming and Java language;
- a working Java compiler (\geq JDK 1.5.x);
- a 0.5 peersim release package or a working peersim source tree (you can download it from sourceforge CVS);
- the Java Expression Parser version \geq 2.3.0 (download it from: <http://www.singularsys.com/jep/>);
- (suggested) gnuplot software.

The aim of this tutorial is to be as practical as possible; the goal is to give the reader the basics of peersim usage and the basics about howto write a simple component. This tutorial IS NOT exhaustive at all!

Because of the spiritus of this tutorial, after a brief introduction to basic concepts, we will try to learn peersim
and its basic components using a by example methodology.

2 Introduction to Peersim

2.1 Why peersim

One of the P2P system properties is that they can be extremely large scale (millions of nodes); another issue to deal with, is the high dynamism of such systems: nodes in the network join and leave continuously. Setting up a protocol experiments in a such simulated environment it's not an easy task at all.

Peersim has been developed to cope with these P2P properties and thus to reach extreme scalability and to support dynamism. In addition, the simulator structure is based on components and makes easy to fast prototype a simulation joining together different pluggable building blocks. The term "components" used here has no relation with high level component architectures (e.g.: CORBA, DOM+).

The peersim performances can be reached only assuming some relaxing assumptions about the simulation details. For example, the overhead introduced by the low level communication protocol stack (e.g.: TCP or UDP) is not taken into account because of the huge additional memory and CPU time requirements needed to accomplish this task. Another simplifying assumption is the absence of concurrency: in peersim the simulation is sequential and based on the concept of cycle in which every node can select a neighbor (the neighborhood relation could be defined by a fixed topology or defined by an overlay management protocol such as *Newscast*) and perform a protocol defined function.

2.2 Peersim simulation life-cycle

The peersim structure is aimed to promote modular programming of building blocks. Every such block is easily replaceable by another component having a similar function, that means, in brief, having the same interface. In the peersim framework, a simulation is carried by the *Simulator* class. The general idea of the simulation model is:

1. choose a network size (number of nodes);
2. choose 1 or more protocol to experiment with and eventually initialize the protocol(s); this step will build a topology on top of raw nodes inserted at the previous point;
3. choose 1 or more *Control*¹ object to monitor the properties you are interested in and to modify or perturb some parameter during the simulation execution (e.g.: the size of the network, update particular values inside protocols, ...);

¹The *Control* is the unified interface that substitutes the old *Observer* and *Dynamics* interfaces.

<i>Node</i>	All the elements of a P2P network are called nodes, the interface manages the local view of neighbor, the reference to the protocol, its index identifier inside the topology global array (invisible to protocols)
<i>CDProtocol</i>	A protocol simply defines an operation to be performed at each cycle (only method <code>nextCycle()</code> is defined)
<i>Linkable</i>	A class implementing this interface has access to the underlying network: can access to its local view of neighbor
<i>Control</i>	Is a very general interface to run any kind of code using its <i>execute()</i> method. In general, it is the base interface to observe or modify the simulation
<i>Vector</i>	Is a package of classes aimed to modify and analyze numeric vectors defined by the vector of protocol instances in the overlay. All the protocol instances contained by nodes in the network define a protocol vector. These classes handle protocol vectors as a whole: for example, by observing one of its fields and reporting aggregation statistics over them.

Table 1: Suggested peersim subset of classes or interfaces to know about.

4. ... run your simulation invoking the *Simulator* class

This is a very general model to give the reader an idea to start with, but it can be extremely more complex.

All the object created during the simulation are instances of classes that implements one or more well defined framework interfaces. The main interfaces I suggest you to become familiar with are in the Table 1.

The life-cycle of a peersim simulation is hard coded inside the *peersim.Simulator* class. It first reads a particular configuration file (see section 2.3) containing all the simulation parameters concerning all the objects involved in the experiment. If no error occurs, the simulator loads and executes *Control* type objects. A special *Control* object (*peersim.cdsim.FullNextCycle*)

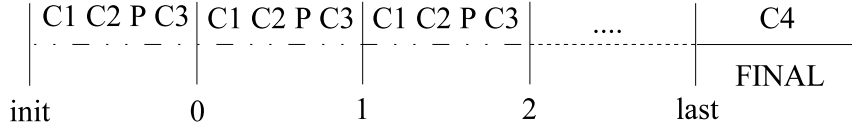


Figure 1: Controls and protocols scheduling. The “C” letters indicate control component, while letter “P” indicates a generic protocol execution. The numbers in lower part of the picture indicates the peersim cycles. After the last cycle, it is possible to run a final control to retrieve a final snapshot.

is dedicated to actually run the protocols From the developer point of view, it is important to note that the protocols creation process is based on **cloning**: only one instance of each protocol is actually forged (with the **new** statement) and then it is cloned to populate all the network. Thus the *clone()* method has to be designed with care to avoid unpredictable results.

Each object in peersim (controls and protocols) is wrapped into *Scheduler* objects which adds fine grained scheduling facilities to each simulation component. This approach leads to total freedom in the cycle execution. For example, it is possible to run a subset of protocols at the beginning, then some controls in between and finally the other protocols (o even a subset). This scenario is depicted in Figure 1.

The initialization phase is carried out by special *Control* objects that run only at the beginning. To obtain this effect, this initializer is internally wrapped by the simulator in a *Scheduler* class object that ensures a single shot. In the configuration file, the initialization components are easily recognizable by the `init` prefix. Please note that in the following pages we will talk about *Initializer* objects just to remark their function and to distinguish them from ordinary *Control* objects, however there is no *Initializer* class in peersim ².

For more *Control* objects after the *Dynamics* object(s).

When a *Control* has to collect data, they are formatted and sent to standard output and can be easily redirected to a file to be collected for further work. When developing such a *Control* class, a good rule is to print out data using the *peersim.core.Log* static class.

2.3 The config file

The config file is a plain ASCII text file, basically composed of key-value pairs; the lines starting with “#” character are ignored (comments). The pairs are collected by a standard Java *java.util.Properties* object when the simulator starts using for example the following command:

²Class *Initializer* was available in early peersim releases, but since version 0.3 has been completely removed.

```
java -cp <class-path> peersim.Simulator config-file.txt
```

Clearly the classpath is mandatory only if you have not set it yet in a global shell variable.

2.4 Configuration example 1

First of all, what we are going to do in this first experiment?

We are going to impose a fixed P2P random topology composed by 50000 node network; the chosen protocol is *Aggregation* (what is aggregation? see appendix A) using an average function. The values to be aggregated (averaged) at each node are initialized using a linear distribution on the interval [0, 100]. Finally a *Control* monitors the averaging values. Looks easy!

```
1 # PEERSIM EXAMPLE 1
2 # random.seed 1234567890
3 simulation.cycles 30
4 control.0 peersim.cdsim.Shuffle
5
6 overlay.size 50000
7 overlay.maxsize 100000
8
9 protocol.0 peersim.core.IdleProtocol
10 protocol.0.degree 20
11
12 protocol.1 example.aggregation.AverageFunction
13 protocol.1.linkable 0
14
15 init.0 peersim.dynamics.WireKOut
16 init.0.protocol 0
17 init.0.k 20
18
19 init.1 peersim.vector.LinearDistribution
20 init.1.protocol 1
21 init.1.max 100
22 init.1.min 1
23
24 control.1 example.aggregation.AverageObserver
25 control.1.protocol 1
```

Lets comment the code line by line. The first thing to note are the key names: some of them are indexed and some other not (e.g. protocol.0.xxx versus simulation.<parameter>). That means the un-indexed keys refer to static simulation elements, in fact the simulation itself is one (is a singleton) and the same holds for the P2P network: only one network!

Please note that from Peersim release 0.2 the component indexes in the configuration are no more restricted to numbers, but can be any user-readable string.

For the other simulation components you can think about the existence of a dedicated array for each of their type (one for protocols, initializer, ...); the index is the only reference to deal with them. So the key for indexed components can be (informally) expressed as:

```
<init | protocol | observer | dynamics> . index_number [. <parameter_name> ]
```

The final `<parameter_name>` is contained between `[]` to express that it is optional. This is the case when the element is declared. For example, at **line 9**, the first protocol chosen comes to life; the **key part** contains its type (or interface type) followed by the index and **the value part** contains the desired component class full package path (you have to check the javadoc files or the source tree to discover the correct package path). In the case of a component parameter declaration, the **key part** contains the parameter name and the **value part** is simply the value desired (usually an integer or a float).

At this point, should be clear that **from line 3 to line 7** some global simulation properties are imposed; these are the total number of simulation cycles and the overlay network size. The *Shuffle* control (**line 4**) is needed to shuffle the order in which the nodes are visited in each cycle. The parameter `overlay.maxsize` (**line 7**) sets an upper bound on the network size, but in this example it is useless (you can comment it out) and it is only present for sake of completeness (it will be useful next).

From **line 9 to line 13**, two protocols are put in the arena. The first one, *peersim.core.IdleProtocol* does nothing. It is useful because of its ability to access to the topology, in fact it provides neighbour links to each node. This feature is present because *IdleProtocol* is an implementation of the *Linkable* interface. Next line declares the graph degree.

The second protocol (`protocol.1 aggregation.AverageFunction`) is the averaging flavor of aggregation. Its parameter (`linkable`) is extremely important: it expresses the need to access the topology using not this protocol itself (aggregation), but with the linkable-implementing underlying protocol (the one the parameter index is referencing to). This is due to the structure of aggregation: it does not implement the *Linkable* interface, so it can not “see” the neighbor list by itself and it must use some other protocol in order to do that. The value of parameter `linkable` is the index of a *Linkable* interface implementing protocol (*IdleProtocol* in the example). Clearly to know if a protocol can get access to the topology directly or not, you have to check the documentation (or source code).

From **line 15 to line 22**, it is time to initialize all the components previously declared. Again, the initialization components are 2 and are indexed as usual. The first initializer `peersim.init.WireKOut`, imposes a node wiring (e.g., a topology). The nodes using the declared protocol are

linked randomly to each-other to form a random graph having the specified degree (**k**) parameter. Please note that this degree declaration is exactly the same of the previous (the one dedicated to the first protocol creation).

The second initializer task is to initialize the aggregation function value-field to be averaged. The initialization values follows a linear distribution fashion. The parameters declared are three: **protocol**, **max**, **min**. Respectively, their meaning is:

- a protocol to point to: the initializer needs a reference (or index, as in the provided example) to a protocol extending *aggregation.AbstractFunction* Class to get access to the value to be aggregated (averaged); it is clear that this protocol must be *aggregation.AverageFunction* (index 1);
- the maximum value in the linear distribution;
- the minimum value in the linear distribution

Finally at **line 24,25** the last component is declared: **aggregation.averageObserver**. Its only parameter used is **protocol** and clearly refers to the *aggregation.AverageFunction* protocol type, so the parameter value is index 1 (in fact: **protocol.1 aggregation.AverageFunction**).

Now you can try the example writing on a console the following line:

```
java -cp <class-path> peersim.Simulator example1.txt
```

The class-path is mandatory only if the current system has not peersim classes in the shell CLASSPATH environment variable. To get the exact output that will follow, the reader should uncomment the parameter at **line 2**:

```
random.seed 1234567890
```

on top of the configuration file. This parameter is very useful to replicate exactly the experiment results based on (pseudo) random behavior. The experiment output is (some initialization strings may be different):

NOTE: cut and paste the experiment:

```
Simulator: loading configuration
ConfigProperties: File example/config-example1.txt loaded.
Simulator: starting experiment 0 invoking peersim.cdsim.CDSimulator
CDSimulator: resetting
Network: no node defined, using GeneralNode
CDSimulator: running initializers
- Running initializer init.0: class peersim.dynamics.WireKOut
- Running initializer init.1: class peersim.vector.LinearDistribution
CDSimulator: loaded controls [control.0, control.1]
CDSimulator: starting simulation
observer.0 0 28.57969570575493 1.0 50.49999999999998 100.0 1.0 50000 50000
```

```

Simulator: cycle 0 done
observer.0 1 15.744375112466432 0.5508937280006126 50.5000000000000185 99.64260285205704 1.993979879597592 50000 50000
Simulator: cycle 1 done
observer.0 2 8.77307045667709 0.3069686446980087 50.500000000000009 86.06868887377748 11.048700974019479 50000 50000
Simulator: cycle 2 done
observer.0 3 4.909681896225926 0.17178915922597776 50.499999999999794 74.03587220181905 22.769780085543115 50000 50000
Simulator: cycle 3 done
observer.0 4 2.7403309556342257 0.09588383948687113 50.5000000000000426 65.43171163227953 33.427798365537626 50000 50000
Simulator: cycle 4 done
observer.0 5 1.538286672869342 0.053824459459153234 50.49999999999973 59.82515640226745 42.62594413722992 50000 50000
Simulator: cycle 5 done
observer.0 6 0.866397905938638 0.03031515502679675 50.500000000000007 55.26130498088358 45.94325388089578 50000 50000
Simulator: cycle 6 done
observer.0 7 0.485544546348093 0.016989143318636584 50.49999999999996 53.34350686753126 47.92146780934889 50000 50000
Simulator: cycle 7 done
observer.0 8 0.27325943590085566 0.009561313693267594 50.499999999999936 51.953084686348944 49.100818456230826 50000 50000
Simulator: cycle 8 done
observer.0 9 0.15407802503043988 0.005391170942362905 50.499999999999954 51.464657035213264 49.43879802069546 50000 50000
Simulator: cycle 9 done
observer.0 10 0.08620333588583261 0.0030162440067013846 50.5000000000000156 51.099961126584006 49.98131655222747 50000 50000
Simulator: cycle 10 done
observer.0 11 0.04848730705794467 0.0016965648464962858 50.49999999999997 50.816956855036466 50.22577832539035 50000 50000
Simulator: cycle 11 done
observer.0 12 0.027214744249562235 9.522405182250473E-4 50.499999999999524 50.65301253758219 50.29955826845794 50000 50000
Simulator: cycle 12 done
observer.0 13 0.015246845383671713 5.334852246380476E-4 50.500000000000032 50.59479421528527 50.38736504947625 50000 50000
Simulator: cycle 13 done
observer.0 14 0.008587160488627146 3.004636780264248E-4 50.499999999999815 50.543660258997136 50.44122418780829 50000 50000
Simulator: cycle 14 done
observer.0 15 0.004850437249792671 1.697161964119851E-4 50.500000000000037 50.52544970665122 50.46753516145482 50000 50000
Simulator: cycle 15 done
observer.0 16 0.0027428141606463717 9.59707265215587E-5 50.500000000000047 50.515509548126744 50.48203242786418 50000 50000
Simulator: cycle 16 done
observer.0 17 0.001550607390364058 5.4255559832703925E-5 50.49999999999997 50.50982384430303 50.49003444731606 50000 50000
Simulator: cycle 17 done
observer.0 18 8.746858998689715E-4 3.0605150904137896E-5 50.50000000000003 50.50564226243819 50.495105203016905 50000 50000
Simulator: cycle 18 done

```

The observer component produces many numbers, but looking at the 6th and 7th data columns (respectively the maximum of averages and the minimum of averages) it is easy to see how the variance decreases very quickly. At cycle 12 (look at the underlined data), quite all the nodes has a very good approximation of the real average (50). Try to experiment with different numbers and then to change the init distribution (e.g.: using `aggregation.Peak DistributionInitializer`) and / or the protocol stack (put *Newscast* or *SCAMP* instead of *IdleProtocol*).

2.5 Configuration example 2

This second example is an improved version of the first one. What is new? Now the aggregation protocol runs on top of *Newscast* and a new values distribution is provided (simply uncomment it and comment out the first two lines of the previous distribution). Moreover, there is a *Control* object that changes the network size, shrinking it by cutting out 500 nodes each time).

```

1 simulation.cycles 30
2 control.0 peersim.cdsim.Shuffle
3
4 overlay.size 50000
5 overlay.maxsize 200000
6
7 protocol.0 example.newscast.SimpleNewscast

```



```

8 protocol.0.cache 20
9
10 protocol.1 example.aggregation.AverageFunction
11 protocol.1.linkable 0
12
13 init.0 peersim.dynamics.WireKOut
14 init.0.protocol 0
15 init.0.k 20
16
17 #init.1 example.aggregation.PeakDistributionInitializer
18 #init.1.value 1
19 init.1 peersim.vector.LinearDistribution
20 init.1.protocol 1
21 init.1.max 100
22 init.1.min 1
23
24 control.1 example.aggregation.AverageObserver
25 control.1.protocol 1
26
27 control.2 peersim.dynamics.DynamicNetwork
28 control.2.add -500
29 control.2.minsize 4000
30 control.2.from 5
31 control.2.until 10

```

The global parameters are the same as in the previous example; only new additions are discussed below. At **line 7-8** there is the *Newscast* (what is newscast? see Appendix B) component declaration with its only parameter `cache` (please note: cache size should be at least as large as network `k` degree size). At **line 17-18** there is a different distribution type: `aggregation.PeakDistributionInitializer`, but it is inactive. To switch it on, simply delete the preceding symbol “#” and comment out the following 4 lines. The peak distribution initializes all nodes except one with 0 value and the node left takes the value declared in parameter `value`.

From **line 27 to 32** is present the last new component: `control.2 peersim.dynamics.DynamicNetwork`. As stated previously, a *Control* implementing object can be used to to change some other object properties; the change can be performed at each simulation cycle (default behavior) or using a more sophisticated approach. The object chosen in the example deletes 500 nodes from the net at each time (well, it is not completely correct to talk about deletion in peersim vision, since the *Linkable* interface does not support the delete operation; so it’s better to think about “unlinking” nodes from the overlay). The parameters `add`, `minsize`, `from` and `until` have respectively the following meaning:

- adds the specified number of nodes (if negative, subtracts);
- the minimum size is referred to the overlay: it can not be less than what is stated here;

- the cycle number from which the component can start running;
- the cycle number until which the component object can run.

Other parameters are available; please check the source or the JavaDoc.

It is interesting to note that not all the parameters associated to a *Control* component can be found in its source code (or documentation); this is due to the Simulator class behavior. When it creates the *Control* instances, it wraps them in a *Scheduler* class object: this is the class where some parameters (such as **step**, **from**, **until**) are actually defined.

2.6 Advanced configuration features

Thanks to the presence of the Java Expression Parser (since release 0.2), the configuration file can handle many types of **expressions**, such as boolean expressions, common mathematical functions and well known predefined constants (e.g.: π and e); for an exhaustive feature list check the Java Expression Parser web page (<http://www.singularsys.com/jep/index.html>).

Expressions can be used anywhere instead of numeric values, as follows:

```
MAG 2
SIZE 2^MAG
```

the variable **SIZE** will be auto-magically evaluated in number 4.

Multiple expressions can be written in a tree-like fashion and they'll be evaluated recursively (the CPU conscious users have to know that no optimizations are performed and the same expression may be evaluated many times) as in the following code sample:

```
A B+C
B D+E
C E+F
D 1
E F
F 2
```

The evaluation will produce: A=7, B=3, C=4, D=1, E=2 and F=2.

Recursive definitions are not allowed and a simple trick is used to avoid them: if the recursion depth is greater than a configurable threshold parameter (set at 100 by default) an error message is printed and the simulator stops.

For any kind of simulator object (e.g., for any prefix allowed in the configuration file), it is possible to specify an ordering scheme. The default one is given by the “index” alphabetically order. In fact, the object prefixes are not limited to numerical indexes, but they can be expressed by any string, as follows:

```
control.conn ConnectivityObserver
control.0 Class1
control.1 Class2
```

The lexicographical order can be explicitly overridden giving an item name list separated by any non-word character (non alphanumeric or underscore) in the following directive:

```
order.observer 2 conn 0
```

If not all names appear in the list, then the vacant objects will follow the default alphabetical order. For example:

```
order.observer 2
```

will produce the following order:

```
{ observer.2 ; observer.0 ; observer.conn }
```

Another available feature is to tell the simulator which items are allowed to run using the following directive:

```
include.control conn 2
```

This will return `control.conn` and `control.2`. If the list is empty, then an empty ordering array will be generated; means that, in this case, no controls will run.

2.6.1 A concrete example

To have a practical idea about how to use these new features, the following example is presented; it is a modified example2 version.

```
1 #random.seed 1234567890
2 simulation.cycles 30
3 control.0 peersim.cdsim.Shuffle
4
5 # Imposes the correct protocol running order:
6 order.protocol ncast avgagr
7 include.control avgobs
8
9 overlay.size 50000
10 overlay.maxsize 200000
11
12 protocol.ncast example.newscast.SimpleNewscast
13 protocol.ncast.cache 20
14
15 protocol.avgagr example.aggregation.AverageFunction
```

```

16 protocol.avgagr.linkable ncast
17
18 init.wrr peersim.vector.WireKOut
18 init.wrr.protocol ncast
20 init.wrr.k 20
21
22 # UNCOMMENT THE FOLLOWING LINES TO GET A LINEAR DISTRIBUTION
23 init.ldistrib peersim.vector.LinearDistribution
24 init.ldistrib.protocol avgagr
25 init.ldistrib.max 100
26 init.ldistrib.min 1
27
28 control.avgobs example.aggregation.AverageObserver
29 control.avgobs.protocol avgagr
30
31 control.dnetwork peersim.dynamics.DynamicNetwork
32 control.dnetwork.add -500
33 control.dnetwork.minsize 4000
34 control.dnetwork.from 5
35 control.dnetwork.until 10

```

In this configuration file, the component indexes are no more used, but a string identifier is used instead. Because of the chosen protocol symbol names (`ncast` and `avgagr`), it is necessary to impose a different running order scheme to let newscast run first using (at **line 6**):

```
order.protocol ncast avgagr
```

In addition, a similar approach is used with the controls for observing the simulation (at **line 7**):

```
include.control dnetwork
```

This allows only the `dnetwork` control to run. This approach avoids the usage of error prone comments in the file to switch components on or off.

3 Writing a new protocol

This section covers the description of how to write a new protocol.

3.1 Which kind of protocol?

The protocol we are going to develop is a simple load balancing algorithm. It works as follows. The state of a node is composed of two values: the local load and the quota. The second one is the amount of “load” the node is allowed to transfer at each cycle. The quota is necessary in order to make real load balancing, otherwise it would be simply averaging. Every node contacts the most **distant** neighbor in its local view and then exchanges

at maximum the quota value. The concept of “distance” is expressed in terms of maximally different load from the current node load. Comparing the distance to the actual node load, the protocol chooses to perform a load balance step using a push or pull approach.

After each cycle, the quota value is restored to allow further computation. The protocol does not care about topology management and relies on other components to get access to this feature (e.g.: *Newscast* or *IdleProtocol*).

3.2 Needed components

Now we have a general idea on what we want to code and it is time to adapt it to the *peersim* framework. Writing the protocol class itself, it is usually not sufficient. Some companion components are required. For example, to restore the quota value for each node at the end of each cycle, a specific *Control* object is required. *Peersim* is basically a collection of interchangeable components, so the development of new stuff should have **modularity** in mind and should maximize code reuse. To achieve this, the following classes are proposed:

- **protocol class itself:** it is built on *peersim.vector.SimpleValueHolder*; it is a simple base class to access a single float variable. It shares the same interface as aggregation: many other components can be used together with the load balancing protocol, such as the initializers classes.
- **Control components:**
 - **ResetQuota:** it is necessary to restore the quota value at each node at the end of each cycle (as previously stated). This object is quite straightforward: it simply implements the only one method the interface *Control* declares, invoking the protected protocol method *resetQuota()*
 - **QuotaObserver:** a control to monitor the **quota** parameter and thus the amount of traffic exchanged in the overlay.
 - **Initializer components:** they are not really needed! In fact the aggregation initializers can be used directly because they share the same interface (both extends *SingleValueHolder*). Please note that the initializers provided in the example package are “light”, demo versions; the developer is encouraged to use the *peersim.vector.** package initializers.
 - **Observer components:** the aggregation observers can be used (the *aggregation.AverageObserver* in particular) since they share the same interface.

To give the reader an idea about the actual code to write, the following subsections present code with comments and other deeper explanations.

3.3 The core load balancing class

```
package example.loadbalance;
import peersim.config.Configuration;
import peersim.config.FastConfig;
import peersim.core.*;
import peersim.vector.SingleValueHolder;
import peersim.cdsim.CDProtocol;

public class BasicBalance extends SingleValueHolder implements CDProtocol{
// Fields:
public static final String PAR_QUOTA = "quota"; // allowed config file parameter
private final double quota_value; // original quota value taken from configuration

protected double quota; // current cycle quota

// Constructor:
public BasicBalance(String prefix, Object obj) {
    super(prefix);
    // get quota value from the config file. Default 1.
    quota_value = (double)(Configuration.getInt(prefix+"."+PAR_QUOTA, 1));
    quota = quota_value;
}
```

It is simply standard Java code until now; the class needs also to implement *peersim.cdsim.CDProtocol* (and *Protocol*) interface(s) and to provide the *nextCycle()* method that is where the actual protocol algorithm is located. In addition, the protocols extends the *SingleValueHolder* class, an implementation of *SingleValue* interface. It is a simple solution to have a public standard access (getter and setter methods) to a single internal variable. In this example the variable holds the node actual load.

In the constructor signature, two parameters are present. The first one is a string corresponding to the configuration file protocol key (e.g.: `protocol.1` in the *LoadBalance* protocol case), the second one is the own protocol id index casted in a *Object* type.

```
// Resets the quota.
protected void resetQuota() {
    this.quota = quota_value;
}
```

The *resetQuota()* method is called by a special control object at the cycle end. Clearly a suitable control entry should be present in the configuration file (such as: `control.0 loadbalance.ResetQuota` and `control.0.protocol protocol-index`). This method is not mandatory, but it is much more software engineering oriented than a dirty variable access performed by the

dynamics object.

```
public Object clone() throws CloneNotSupportedException {
    BasicBalance af = (BasicBalance)super.clone();
    return af;
}

// Implements CDProtocol interface
public void nextCycle( Node node, int protocolID ) {
    int linkableID = FastConfig.getLinkable(protocolID);
    Linkable linkable = (Linkable) node.getProtocol( linkableID );
    if (this.quota == 0) {
        return; // skip this node
    }
    // this takes the most distant neighbor based on local load
    BasicBalance neighbor = null;
    double maxdiff = 0;
    for(int i = 0; i < linkable.degree() ; ++i)
    {
        Node peer = linkable.getNeighbor(i);
        // The selected peer could be inactive
        if(!peer.isUp()) continue;
        BasicBalance n = (BasicBalance)peer.getProtocol(protocolID);
        if(n.quota!=1.0) continue;
        double d = Math.abs(value-n.value);
        if( d > maxdiff ) {
            neighbor = n;
            maxdiff = d;
        }
    }
    if( neighbor == null ) {
        return;
    }
    doTransfer(neighbor);
}
```

The first method is required by the *Protocol* interface and basically calls the ancestor cloning method. So, nothing special here.

The second one is required by *CDProtocol* interface. It is the behavior performed by the protocol. The arguments represent a reference to the node itself (the node on which the simulator is invoking the *nextCycle()* method) and the index protocol identifier (the *BasicBalance* protocol id in this case). First it has to get a reference (in indexed form) to the *Linkable* interface enabled protocol in the node protocol stack; as a remind, something implementing the *Linkable* interface, is an entity capable of accessing the topology. Having this linkable reference we can access to the real *Linkable* interface implementation with:

```
int linkableID = FastConfig.getLinkable(protocolID);
```

```
Linkable linkable = (Linkable)node.getProtocol(linkableID);
```

Using the static *peersim.config.FastConfig* class we can get the current protocol corresponding Linkable identifier; this class manages the protocol *linkable* parameter without direct user intervention. Then we can access the actual linkable object as shown in the second line.

If the local quota is equal to 0, means that the node we have already spent its amount of network traffic, so it returns.

To get the most distant node from the current one, a for loops on all neighbor node load value; the number of neighbor is equal to the node degree (accessible thanks to *Linkable* interface). To pick a node having a the *Linkable* access:

```
Node peer = linkable.getNeighbor(i);
```

and from this obtained *Node* interface reference it is possible to get the protocol interface we are interested in (*BasicBalance*):

```
BasicBalance n = (BasicBalance)peer.getProtocol(protocolID);
```

When the protocol finds a suitable neighbor, it performs a load balancing step invoking the *doTransfer()* method.

```
// Performs the actual load exchange selecting to make a PUSH or PULL approach.
// It affects the involved nodes quota.
protected void doTransfer(BasicBalance neighbor) {
    double a1 = this.value;
    double a2 = neighbor.value;
    double maxTrans = Math.abs((a1-a2)/2);
    double trans = Math.min(maxTrans,quota);
    trans = Math.min(trans, neighbor.quota);

    if( a1 <= a2 ) // PULL
    {
        a1+=trans;
        a2-=trans;
    }
    else // PUSH
    {
        a1-=trans;
        a2+=trans;
    }

    this.value = a1;
    this.quota -= trans;
    neighbor.value = a2;
    neighbor.quota -= trans;
}
```


The last method takes, as parameter, a reference to the picked neighbor. This is the place where it is time to decide to perform a pull or a push load balancing approach. To make this choice the local load value is compared with the neighbor load value. In case of a push choice, the local value is increased and the other node value is decreased; in the other case (pull) the exact opposite holds. The *maxTrans* variable is the absolute amount of “load” to transfer to reach the balance between the two involved nodes; because of the quota upper bound on the transfers at each cycle, the algorithm chooses the minimum between the quota itself and the aimed *maxTrans* amount. The quota value is decreased by the same amount at both nodes.

3.4 Load balancing control class code

```
package loadbalance;

import peersim.config.*;
import peersim.core.*;

public class ResetQuota implements Control {
    // Fields:
    public static final String PAR_VALUE = "value";

    public static final String PAR_PROT = "protocol";

    private final double value;
    private final int protocolID;

    // Constructor:
    public ResetQuota(String prefix)
    {
        value = Configuration.getDouble(prefix+"."+PAR_VALUE);
        protocolID = Configuration.getPid(prefix+"."+PAR_PROT);
    }

    // Dynamics interface method:
    public void execute() {
        for(int i=0; i < Network.size(); ++i)
        {
            ((BasicBalance)Network.get(i).getProtocol(protocolID)).resetQuota();
        }
    }
}
```

The code is very compact because the *Control* interface itself is very simple: only the *execute()* method. The constructor takes care of initializing the configuration file parameters (respectively: the reset value and the protocol identifier to deal with). The *execute()* method makes use of network global knowledge: it invokes the *resetQuota()* method on all the *Network* object elements (it is a static object available everywhere in the simulator environ-

ment; you can think about it as an array). It is clear that the simulator has global knowledge, but it is up to the protocol developer to make use or not of this facility according to the consistency of the simulation itself.

3.5 Implementing the Linkable interface

In this howto there are a lot of references about the *Linkable* interface and about its importance, so for the sake of completeness, it is time to give a look at how to implement it in brief. It is interesting to note that this interface should be implemented by low level or by topology management protocols and not by a higher level protocol such as a load balancing one. The reason to discourage the implementation in higher level components, is the risk to affect modularity. At least, the reader should consider the ability to switch off the built in *Linkable* interface and to use an external protocol facility instead.

The interface defines five methods: *degree()*, *getNeighbor()*, *addNeighbor()*, *contains()*, *pack()*. These methods are not usually invoked by the protocol itself (except for *getNeighbor()*), but by an *Initializer* object instead (such as: *peersim.vector.WireKOut*). Please note that there is no way to remove nodes from the overlay; the only chance to get a similar effect, is to disable a peer accessing to the *peersim.core.Fallible* interface (extended by the *Node* interface) and setting one of the available node states (*peersim.core.Fallible.OK — DEAD — MALICIOUS — DOWN*).

A feasible implementation could be the following. First of all, the class (e.g.: *BasicBalance*) needs a structure to represent the neighbor view: an *ArrayList* structure is fine.

```
protected ArrayList nView = null;
// Constructor:

public BasicBalance(String prefix, Object obj) {
    ...
    nView = new ArrayList();
    ...
}
// Linkable interface implementation methods:
public int degree() {
    return nView.size();
}

public Node getNeighbor(int i) {
    return (Node)nView.get(i);
}

public boolean addNeighbor(Node n) {
    if (!contains(n)) {
        nView.add(n);
        return true;
    }
}
```

```

    }
    else {return false;}
}

public boolean contains(Node n) {
    return nView.contains(n);
}

public void pack() { ; } // unused!

```

Again the code is quite straightforward. All the elements inside the view are *Node* class (interface) types. All methods are simple functions built upon the *ArrayList* structure. The last method is included in the interface description with the aim to provide a view size compression facility, but it is usually not implemented (the size of each view is typically quite small).

4 A second new protocol

This new protocol is an extensions of the previous one. The general core is quite the same, but the algorithm uses the global load average value instead of the most distant neighbor load value. To calculate the global load average, a little trick is used; it would be possible to calculate this value using aggregation, but we can **simulate** the aggregation effect (alias calculating the average load) by running a static method with global knowledge once. This method will initialize a global variable available to all nodes.

This protocol is targeted to gain advantage from the newscast protocol features; when a node reaches the global load value (average), it switches to a DOWN state. In this way, the node exits from the overlay and the newscast protocol no more cares about it. The effect is that the topology shrinks as soon as the nodes reach the average load.

```

package loadbalance;

import peersim.core.*;
import peersim.config.FastConfig;

public class AvgBalance extends BasicBalance {

    public static double average = 0.0;
    public static boolean avg_done = false;

    // Costructor:
    public AvgBalance(String prefix, Object obj) {
        super(prefix, obj);
    }

    // Method to simulate average function aggregation
    private static void calculateAVG(int protocolID) {
        int len = Network.size();
    }
}

```

```

double sum = 0.0;
for (int i = 0; i < len; i++) {
    AvgBalance protocol = (AvgBalance)Network.get(i).getProtocol(protocolID);
    double value = protocol.getValue();
    sum += value;
}
average = sum / len;
avg_done = true;
}

```

The first part is straightforward. Two global variables are defined: *average*, and *avg_done*; the second is a flag used to be sure not to perform the calculation more than once. A different and much more elegant approach is to define the average calculation method inside a static constructor, but this solution is **wrong**! When the node protocol objects are created, the load distribution is not defined yet, so the global average result will be 0.

The function *calculateAVG()* simulates the average aggregation behaviour. It makes use of global knowledge, looping on each overlay node.

```

protected static void suspend( Node node ) {
    node.setFailState(Fallible.DOWN); }
}

```

This is the utility function to exit from the topology; simply sets a node state from *Fallible* interface.

```

// CDProtocol Interface implementation. Overrides the BasicBalance implementation:
public void nextCycle( Node node, int protocolID ) {
    // Do that only once
    if (avg_done == false) {
        calculateAVG(protocolID);
        System.out.println("AVG only once "+average);
    }

    if( Math.abs(value-average) < 1 ) {
        AvgBalance.suspend(node); // switch off node
        return;
    }

    if (quota == 0 ) return;

    Node n = null;
    if (value < average ) {
        n = getOverloadedPeer(node, protocolID);
        if (n != null) { doTransfer((AvgBalance)n.getProtocol(protocolID)); }
    }
    else {
        n = getUnderloadedPeer(node, protocolID);
        if (n != null) { doTransfer((AvgBalance)n.getProtocol(protocolID)); }
    }
}

```

```

        if( Math.abs(value-average) < 1 ) AvgBalance.suspend(node);
        if (n != null) {
            if( Math.abs(((AvgBalance)n.getProtocol(protocolID)).value-average) < 1 )
                AvgBalance.suspend(n);
        }
    }
}

```

Method *nextCycle()* is the protocol algorithm core. It first checks for the average calculation: if the flag is not set, it performs the computation.

If the difference between the current and the average load is less than 1 (the fixed quota value per cycle) the node is suspended and thus exits from the topology defined by the newscast protocol; moreover, if the quota has been already spent, it returns. The protocol then checks if the local value is less or greater than the average and respectively get the most loaded or the least loaded neighbor and exchange.

```

private Node getOverloadedPeer(Node node, int protocolID) {
    int linkableID = FastConfig.getLinkable(protocolID);
    Linkable linkable = (Linkable) node.getProtocol( linkableID );

    AvgBalance neighbor=null;
    Node neighborNode = null;
    double maxdiff = 0.0;
    for(int i = 0; i < linkable.degree(); ++i) {
        Node peer = linkable.getNeighbor(i);
        if(!peer.isUp()) continue;
        AvgBalance n = (AvgBalance)peer.getProtocol(protocolID);
        if(n.quota==0) continue;
        if(value >= average && n.value >= average) continue;
        if(value <= average && n.value <= average) continue;
        double d = Math.abs(value-n.value);
        if( d > maxdiff ) {
            neighbor = n;
            neighborNode = peer;
            maxdiff = d;
        }
    }
    return neighborNode;
}

private Node getUnderloadedPeer(Node node, int protocolID) {
    int linkableID = FastConfig.getLinkable(protocolID);
    Linkable linkable = (Linkable) node.getProtocol( linkableID );

    AvgBalance neighbor=null;
    Node neighborNode = null;
    double maxdiff = 0.0;
    for(int i = 0; i < linkable.degree(); ++i) {
        Node peer = linkable.getNeighbor(i);
        if(!peer.isUp()) continue;
        AvgBalance n = (AvgBalance)peer.getProtocol(protocolID);
        if(n.quota==0) continue;

```

```

        if(value >= average && n.value >= average) continue;
        if(value <= average && n.value <= average) continue;
        double d = Math.abs(value-n.value);
        if( d < maxdiff ) {
            neighbor = n;
            neighborNode = peer;
            maxdiff = d;
        }
    }
    return neighborNode;
}
} // end of class

```

The methods to get the the most loaded or the least loaded neighbor are straightforward and very similar, but are shown for completeness.

5 Evaluating the protocols

The performance about load variance reduction can be analyzed with an *aggregation.AverageObserver* or a *loadbalance.LBObserver* (they are very similar), but do not expect huge differences. In fact, from this point of view, the two protocols have nearly an identical performance, no matter whatever distribution you are using. The *AVGBalance* protocol improvement over the *BasicBalance* one is about the achieved overall load transfer. The *AVGBalance* amount of transfer is minimal and it is practically the same of the theoretical minimal amount of transfer needed to solve the problem (more about this: <http://www.cs.unibo.it/bison/publications/modular-p2p.pdf>).

The *Control* class code to inspect the load transfer amount is the following.

```

package loadbalance;

import peersim.core.*;
import peersim.config.*;
import peersim.util.*;

public class QuotaObserver implements Control {

    public static final String PAR_PROT = "protocol";
    private final String name;
    private final int pid;    // protocol id to monitor
    private IncrementalStats stats;    // object to compute statistics

    // Constructor:
    public QuotaObserver(String name) {
        this.name = name;
        pid = Configuration.getPid(name + "." + PAR_PROT);
        stats = new IncrementalStats();
    }
}

```

```

}

// Observer interface implementation:
public boolean execute() {
    for (int i = 0; i < Network.size(); i++) {
        BasicBalance protocol = (BasicBalance) Network.get(i).getProtocol(pid);
        stats.add( protocol.quota );
    }
    System.out.println(name+": "+stats);
    return false;
}

} // end of class

```

The idea is very simple: at each simulation cycle it collects all node values about quota and prints statistics on the console.

A A few words about Aggregation

It is a very fast epidemic-style protocol targeted to compute a particular function (e.g.: average, max, min, ...) on a numeric value holded at each network node. In order to work, every node needs access to its neighbor list view on the overlay network; no particular requirements about the topology management protocol are imposed. In the case of averaging function, a generic method *updateState(a, b)* returns $(a+b)/2$, where a and b are values at (respectively) node a and node b. This kind of computation is performed by each node at each simulation cycle.

The global average is not affected, but the variance over all the estimates decreases very fast in a few cycles (the convergence rate is exponential and it does not care about the network size). The aggregation protocol is also very robust in case of node failures.

Suggested readings: project BISON publication page <http://www.cs.unibo.it/bison/pub.shtml>.

B A few words about Newscast

Newscast is an epidemic content distribution and topology management protocol. Every peer in the system has a partial view knowledge about the topology which is modeled as a fixed size (c) set of node descriptors. Each descriptor is a tuple consisting of a peer address and a time-stamp recording the time when the descriptor was created.

Each node updates its state by choosing a random neighbor and exchanging with it the view. The exchanging process merges the two involved peer partial view, keeping the c freshest descriptors. in this manner, old information (descriptor) are auto-magically removed from the system as time

goes on. This process allows the protocol to repair the overlay topology removing dead links with minimum effort and this is a great feature for a highly dynamic oriented system where nodes join and leave continuously.

The protocol relies on the timestamps, but it doesn't need synchronized clocks: timestamps have to be only mutually consistent. To achieve this, a simple time normalization of the received descriptors is performed. So time precision is not critical at all.

The emergent topology from newscast topology management has a very low diameter and it is very close to a random graph with (out) degree c . Suggested readings: "Large-Scale Newscast Computing on the Internet" <http://citeseer.nj.nec.com/jelasity02largescale.html>.

C Major peersim changes from previous release

Peersim release 0.5 represents a major change from the previous version. Due to strong refactoring, backward compatibility is broken. The following is a list of the major changes from the previous release. This list is given in order to provide the reader a fast reference to port previously developed components to the new simulator release.

- No more *Observer* and *Dynamics* interfaces. Use **Control** instead.
- All cycle driven stuff is now in *cdsim* package; the Scheduler is now model independent.
- *CommonState* has now a companion class called **CDState**. The former is common to all simulation components, while the latter is common to all cycle driven components.
- No more *CommonRandom*, now use the static field **r** in *CommonState*.
- No more *peersim.util.Log*, now it is in **peersim.core.Log**.
- No more *simulation.shuffle* parameter, shuffling can now be added by the control **cdsim.Shuffle**.
- The very popular parameter **degree** is renamed in **k**.
- Refactored *Wire** controls. They have a common superclass now and the *NodeInitializer* implementations were separated into separate classes (*RandNI* and *StarNI*).
- No more *peersim.dynamics.WireRegularRandom*, now is `textbfpeer-sim.vector.WireKOut`.
- In all classes that used it, now **undir** and **undirected** are both accepted as equivalent parameters.

- The *Wire** classes now have a public **Graph** field that can be used to initialize arbitrary graphs (see JavaDocs of WireGraph).
- Introduced **MethodInvoker** control. It invokes a void method on a specified protocol.
- Implemented the Tarjan algorithm to detect the strongly connected components in GraphAlgorithms.
- Added **peersim.dynamics.WireRegRootedTree** that can wire a regular rooted tree.
- No more *GrowingNetwork*, now *DynamicNetwork* provides its features.
- Added **peersim.dynamics.WireByMethod**: it takes a Linkable protocol and adds connections using an arbitrary user defined (static) method.